

# EAV Antipattern Database Structure

by

Tom Anderson

# EAV Tables

EAV Splits traditional table attributes into atomic tables

---

Columns

Value	Value	Value
Value	Value	Value
Value	Value	Value
Value	Value	Value

Rows

Attributes

Value	Value	Value
Value	Value	Value
Value	Value	Value
Value	Value	Value

Entity

But attributes must be grouped in order to  
make sense of table data

# EAV Tables

## Attribute Sets Table

---

<b>attribute_sets</b>
attribute_set_key - Primary key
ref_user - The user that owns the attribute set.
name - The name of the attribute set

# EAV Tables

## Attributes Table

---

Attribute tables contain column descriptions

<b>attributes</b>
attribute_key - Primary Key
ref_renderer - What renderer does this attribute use?
name - The name of this attribute
is_required - Is the attribute required?
is_unique - Should values of this attribute be unique across all instances of this attribute?

# EAV Tables

## Renderers Table

---

renderers
renderer_key - Primary key
name - The renderer name such as URL, Date or Text
model - The PHP class to handle rendering a form input element and formatting the value. This is an object dictionary.

And Datatype: one of **datetime**, **decimal**, **int**, **text** or **varchar** stored in each model.

# EAV Tables

## Renderer Model

---

```
<?php
```

```
class Renderer_Integer extends Zend_Form_Element_Text implements Collections_Renderer
{
    public $icon = '13.png';
    public $datatype = 'int';

    public function postCreate() {
        $this->addValidator('Int');
    }

    public function formatValue($value) {
        return $value;
    }

    // Preformat value when editing such as date times > dates
    public function formatEditValue($value) {
        return $value;
    }
}
```

# EAV Tables

## Values Tables

---

<b>items_datetime, items_decimal, items_int, items_text, or items_varchar</b>
value_key
ref_attribute - Which attribute does this value implement?
ref_item - Reference to the entity
value - A column type reflecting the table name

EAV encourages no particular way to store values so many ways have been concocted. Many EAV implementations incorrectly use a single value table with a value column of type text. This limits you and doesn't apply the flexibility of datatypes available in a database. **Magento** got it right and uses five value tables for each entity table and I **mimic their work**.

# EAV Tables

## Items Table

---

Entity tables are named after the data it represents e.g. **Items**

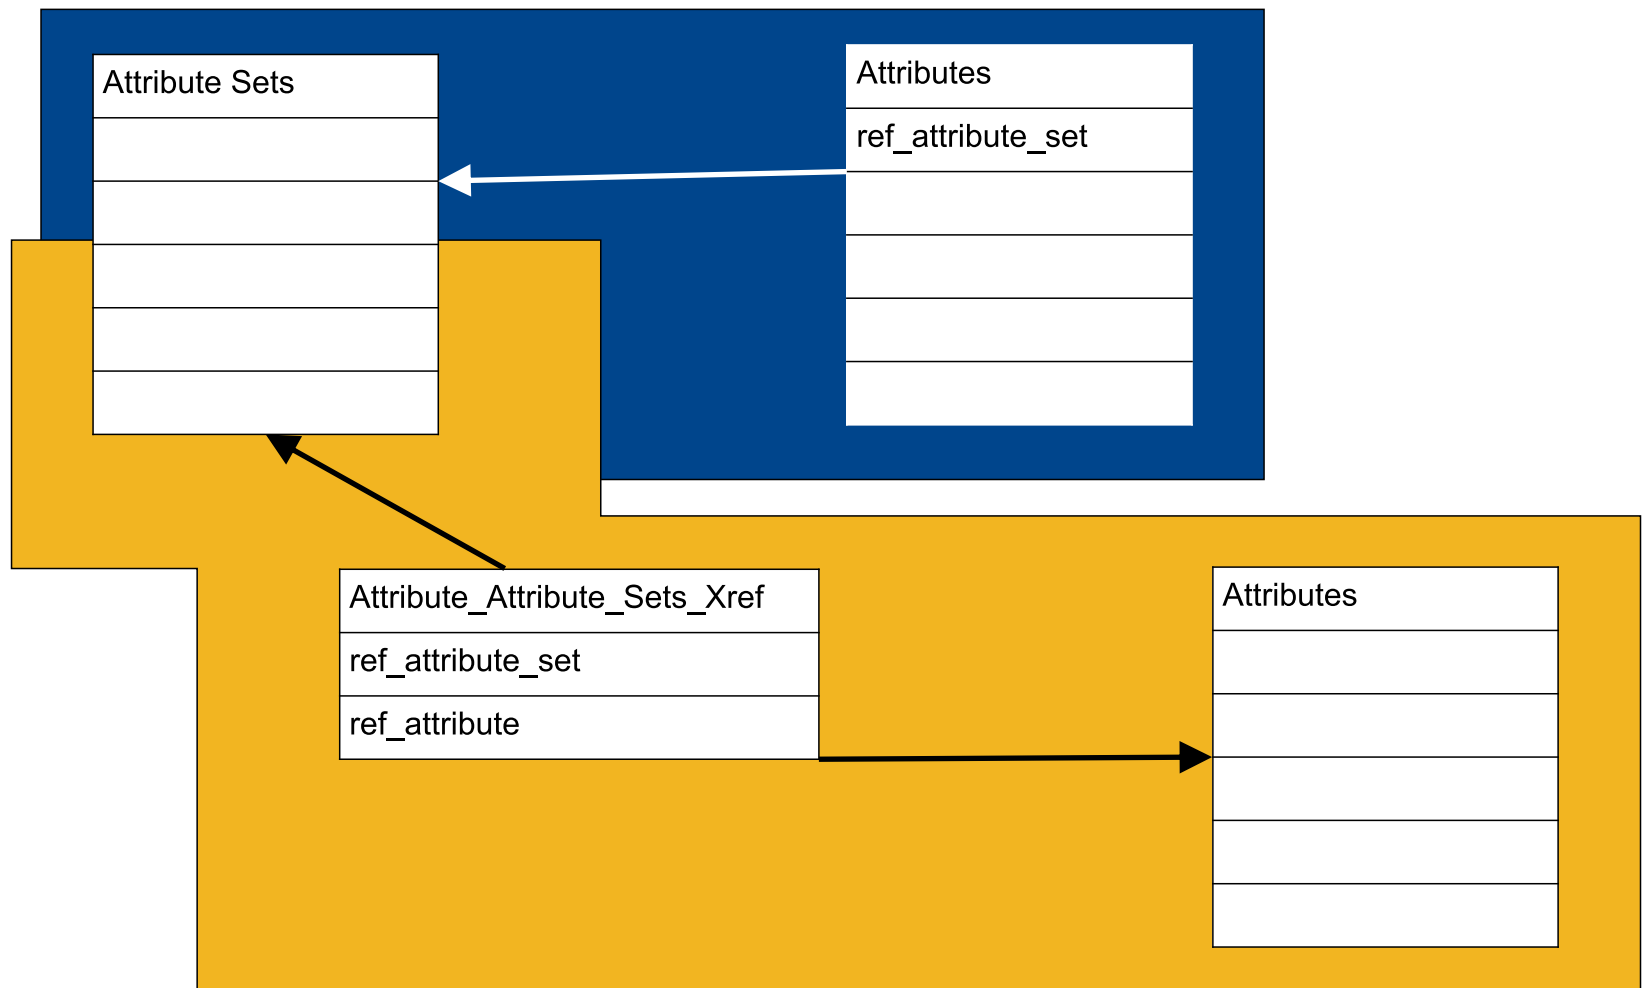
<b>Items</b>
items_key - Primary Key
ref_attribute_set - Which attribute set does this item implement?
ref_item - Does this item have a parent item?
ref_user - The user that created the item.




# EAV Tables

## Attributes - Attributes Sets relationship

Attribute sets may have a many to many or one to many relationship with attributes



Introducing



An MVC & EAV, Domain Model, 100% OO, Collections Engine

# Advantages of Stuki over Omeka

Stuki and Omeka were designed from the same epiphany: that EAV tables could store collections of items with dynamic attributes.

---

## Stuki

Includes Dublin Core Metadata Elements  
Peer Reviewed Edits  
Pictures  
Full Table Auditing  
Watchlists  
Merge Requests  
HTML Purifier  
Ownership Records & User Ownership

**17 Renderers**

**18 Plugins**

### Plugins Include:

Dublin Core COinS  
Favorites  
Footnotes  
Forums  
Intense Debate  
Library of Congress Subject Headings  
PDF Report with URL Barcodes  
Proxy Email  
Social Bookmarks  
Wantlists

## Omeka

Includes Dublin Core Metadata Elements  
Files  
Tags

**31 plugins**

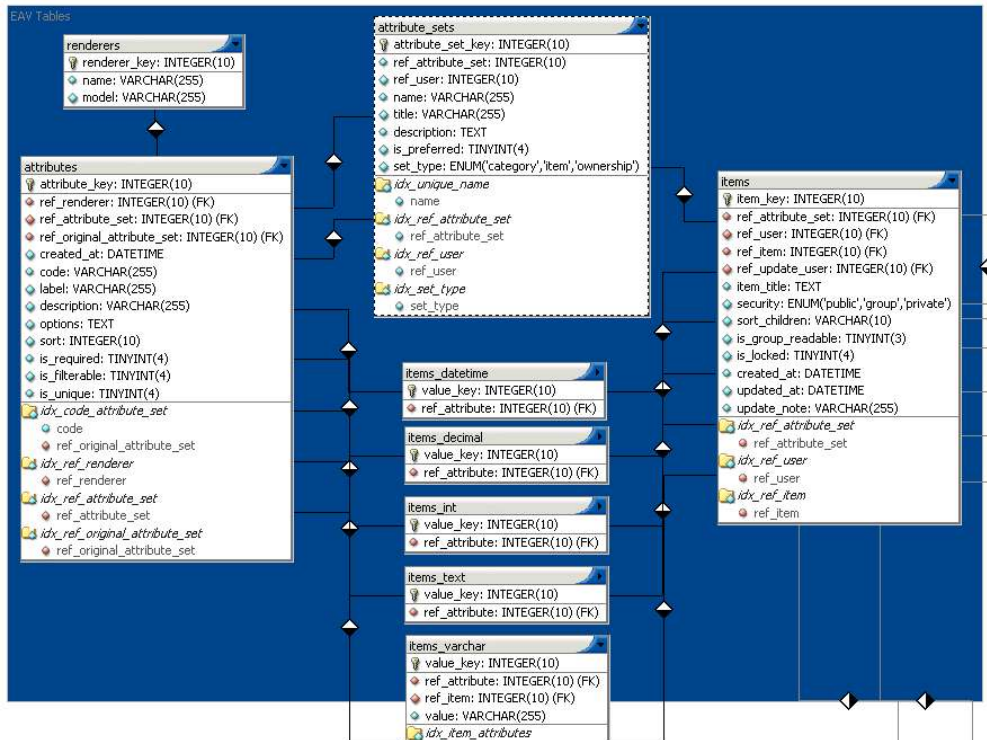
### Including:

Bar Codes and Reports  
Docs Viewer  
Dublin Core COinS  
HTML Purifier  
Intense Debate  
Library of Congress Subject Headings  
Social Bookmarks

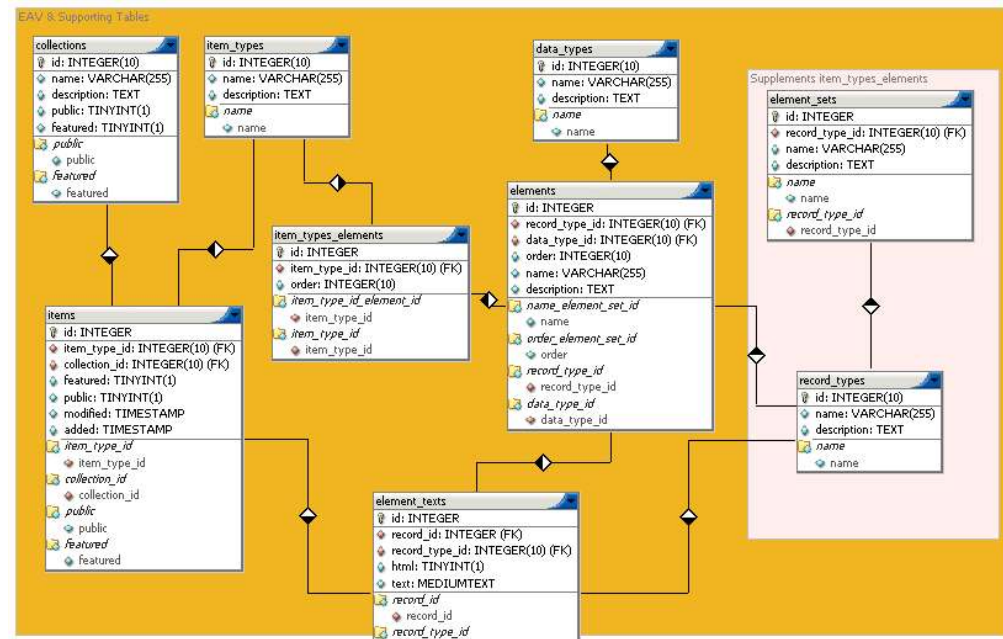
# Advantages of Stuki over Omeka

Stuki has precise, well named EAV tables and values are stored in their appropriate data type following **Magento's** example.

## Stuki



## Omeka



- **attribute\_sets** is similar to **item\_types\_elements**.
- **attributes** is similar to **elements**
- **renderers** is similar to **data\_types** but renderers stores the class each type uses to render it's form element and value. **Omeka** uses **data\_types** values in code throughout the application.
- **item\_types** has no correlation because the 'type' is defined in **attribute\_sets.name** & **description**.
- In **Stuki** a collection is stored like any other item giving them the same flexibility. No **collections** table is necessary.
- **element\_texts.record\_type\_id** duplicates **elements.record\_type\_id**.

- **Items** is similar to **Items** but **Stuki** includes enhanced security, how to sort child elements, a lock, and an update note used in auditing changes.
- MySQL only recognizes one **TIMESTAMP**, not two per table like **items**.
- **element\_texts** is analogous to **items\_datetime**, **items\_decimal**, **items\_int**, **items\_text**, and **items\_varchar**. The significant difference is **Stuki** stores every value in it's appropriate datatype inside the database. **Omeka** stores everything in a **mediumtext** field.
- Please note **element\_texts.record\_id** references **items.id**. This is not represented in the schema and must be divined.
- **element\_sets** and **record\_types** have no analog in **Stuki**. They were added into **Omeka** when other tables were removed.
- In **Stuki** the **attribute** references the **attribute set**. In **Omeka** there is a many to many relationship between **item\_types** and **elements**. **Stuki's** simplistic approach means **attributes** cannot be reused between **attribute sets** but creating new **attributes** is trivial.
- In **Omeka** collections are static objects which must be enhanced with code. In **Stuki** categories, items, regular items, or ownership records are all stored in the **items** table.

# Advantages of Stuki over Omeka

## Searching

### Stuki



#### Scalable, High-Performance Indexing

- \* over 20MB/minute on Pentium M 1.5GHz
- \* small RAM requirements -- only 1MB heap
- \* incremental indexing as fast as batch indexing
- \* index size roughly 20-30% the size of text indexed

#### Powerful, Accurate and Efficient Search Algorithms

- \* ranked searching -- best results returned first
- \* many powerful query types: phrase queries, wildcard queries, proximity queries, range queries and more
- \* fielded searching (e.g., title, author, contents)
- \* date-range searching
- \* sorting by any field
- \* multiple-index searching with merged results
- \* allows simultaneous update and searching

<http://lucene.apache.org/java/docs/features.html>

### Omeka



#### MySQL Full Text Search

Example:

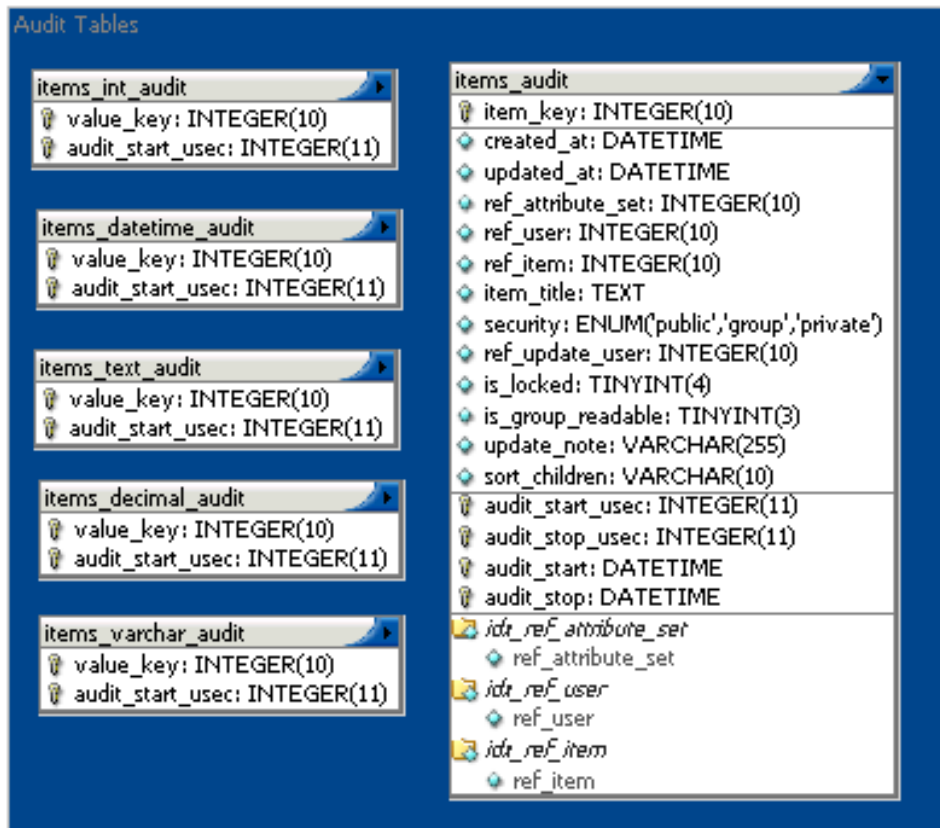
```
SELECT i.id as item_id,  
MATCH (etx.text) AGAINST ('test test') as rank  
FROM ...
```

Note **Omeka** does not use boolean mode thereby reducing the user's options when making searches.

# Advantages of Stuki over Omeka

Stuki has full items auditing extending Joe Celko's approach in SQL For Smarties by adding a microtime stamp to each change.

## Stuki



The screenshot shows a database interface with a left sidebar titled 'Audit Tables' containing five entries: items\_int\_audit, items\_datetime\_audit, items\_text\_audit, items\_decimal\_audit, and items\_varchar\_audit. Each entry shows 'value\_key: INTEGER(10)' and 'audit\_start\_usec: INTEGER(11)'. The main pane shows the 'items\_audit' table with fields: item\_key: INTEGER(10), created\_at: DATETIME, updated\_at: DATETIME, ref\_attribute\_set: INTEGER(10), ref\_user: INTEGER(10), ref\_item: INTEGER(10), item\_title: TEXT, security: ENUM('public','group','private'), ref\_update\_user: INTEGER(10), is\_locked: TINYINT(4), is\_group\_readable: TINYINT(3), update\_note: VARCHAR(255), sort\_children: VARCHAR(10), audit\_start\_usec: INTEGER(11), audit\_stop\_usec: INTEGER(11), audit\_start: DATETIME, and audit\_stop: DATETIME. Below this are three index definitions: idx\_ref\_attribute\_set, idx\_ref\_user, and idx\_ref\_item, each with a single field (ref\_attribute\_set, ref\_user, and ref\_item respectively).

Audit Table	Fields
items_int_audit	value_key: INTEGER(10) audit_start_usec: INTEGER(11)
items_datetime_audit	value_key: INTEGER(10) audit_start_usec: INTEGER(11)
items_text_audit	value_key: INTEGER(10) audit_start_usec: INTEGER(11)
items_decimal_audit	value_key: INTEGER(10) audit_start_usec: INTEGER(11)
items_varchar_audit	value_key: INTEGER(10) audit_start_usec: INTEGER(11)
items_audit	item_key: INTEGER(10) created_at: DATETIME updated_at: DATETIME ref_attribute_set: INTEGER(10) ref_user: INTEGER(10) ref_item: INTEGER(10) item_title: TEXT security: ENUM('public','group','private') ref_update_user: INTEGER(10) is_locked: TINYINT(4) is_group_readable: TINYINT(3) update_note: VARCHAR(255) sort_children: VARCHAR(10) audit_start_usec: INTEGER(11) audit_stop_usec: INTEGER(11) audit_start: DATETIME audit_stop: DATETIME
idx_ref_attribute_set	ref_attribute_set
idx_ref_user	ref_user
idx_ref_item	ref_item

Note the items tables are the only audited tables in Stuki at this time. However, Stuki is crippled at it's current host and triggers cannot be added to all other significant tables in the database. Adding these triggers is a day's job.

## Omeka

No Auditing

# Advantages of Stuki over Omeka

## Final Thoughts

---

Stuki was developed with no knowledge of Omeka.

Stuki has been in development by a solo developer since July 2010. Omeka has been in development for the past two years. Omeka has had a longer period and more eyes for testing and development than Stuki.

The database design and choices made about how to implement the functionality which makes this sort of application applicable are superior in the Stuki. For instance, the choice to create unique attributes for each definition in Stuki is a stronger position than creating attributes and assigning them to definitions, forcing the user to select from a long list of attributes and giving them no preview of what their selection looks like. Stuki considered this approach and decided it was too confusing and complicated the database structure. Stuki, instead, uses unique attributes for each definition thereby making the user choose only the data type and giving them a preview of each as they choose. In turn this streamlined the database structure into a single reference from attributes to attribute types.

Additionally Stuki took it's EAV design from Magento, the undisputed leader in e-commerce applications today. The Magento team chose to store values in five tables so the data types of each could be directly used. This makes it impossible to change a to a different datatype renderer once data exists which uses a definition.

The design of renderers which use specific Zend\_Form\_Element definitions is far superior to Omeka's use of data\_types, the values of which are scattered, hard coded, throughout the application.

In Stuki, every dynamic piece of information is stored in the ini file. There are 2 DEFINE's in Stuki as recommended by Zend. There are no global variables in Stuki. paths.php defines 39 variables essentially used as global variables throughout the application.

Omeka uses global functions thereby breaking a true object oriented approach. Stuki has no globals or global functions and is a true 100% object oriented application.



# Advantages of Stuki over Omeka

## Thoughts on Modeling

---

**Omeka** defines base directories so theoretically you could move your directories to different machines or drives. This is unnecessary in code intended for Unix installations since mapping devices and drives to a directory structure is simplistic. Windows provides analogous pathways.

Labeled “helpers” in **Omeka**, the application loads file after file of global functions(1). This approach to PHP programming lost favour with PHP 4. Fully OO programming has been preferred in every OO language since java's rise to popularity. PHP supports fully OO applications and frameworks like Zend are the preferred approach today. **Omeka** is a hybrid of structural programming with OO libraries.

**Omeka** uses a Domain Model architecture (Active Record). This approach is widely accepted: “...the diffused opinion is that the more complex the business logic and the data involved, the more the application benefits from a rich Domain Model.”(2) I argue the simplistic nature of EAV tables doesn't warrant the use of such a structured model. Leaving each model to make SQL queries as they see fit it imbues flexibility in the core of the application.

**Omeka** takes the Domain Model in directions it's not designed. The ancestor `Omeka_Controller_Action` is assigned a table which that controller controls. This irreparably breaks the MVC design because it is establishing business logic on the controllers. An MVC application should be scriptable from the command line with no use of controllers.

**Stuki** uses business logic models which in turn use each other. The loading of models in **Stuki** is done with `$items = Collections_ModelBroker::get('Items');` The model broker then creates a new `Collections_Model_Wrapper(new $modelName);` This wrapper and it's enclosed model are then placed in static memory so they are only created once (unless a model's flagged property to always recreate is true).

The `Collections_Model_Wrapper` allows plugins to wrap around a specific model function and override the function completely or change parameters sent to or received from the function. The mechanism is tested but I have not created a plugin which takes advantage of this; yet but the overhead is trivial. Plugins are generally satisfied with access to the Smarty template before it's displayed.

Plugins are modelled from `Zend_Controller_Action_Helper` and all plugins are descended from `Zend_Controller_Action_Helper_Abstract`(3). These are not related to **Omeka's** “helpers”. See reference article.

**Stuki** uses and loads models on an as-needed basis. No preloading of 'mixins' such as `Omeka_Record_Mixin` which does pseudo-multiple inheritance. I'm not a proponent of multiple inheritance.

1: `~/application/helpers/all.php`

2: <http://css.dzone.com/books/practical-php-patterns/practical-php-patterns-domain>

3: <http://devzone.zend.com/article/3372>



# EAV Antipattern Database Structure

by

Tom Anderson